

Tuning Hyperparameters Using Cross Validation Scores



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Hyperparameter tuning in model validation

Model parameters vs. model hyperparameters

Hyperparameter tuning using Azure ML Studio

Hyperparameters

Model configuration properties that define a model, and remain constant during the training of the model

Hyperparameters

Can be thought of as part of model design

Understanding Hyperparameters

Model Inputs

Model Parameters

**Model
Hyperparameters**

Understanding Hyperparameters

Model Inputs

Input data points, training
dataset

Model Parameters

Model Hyperparameters

Understanding Hyperparameters

Model Inputs

Input data points, training dataset

Model Parameters

Structure of decision tree

Model Hyperparameters

Understanding Hyperparameters

Model Inputs

Input data points, training dataset

Model Parameters


Structure of decision tree

Model Hyperparameters

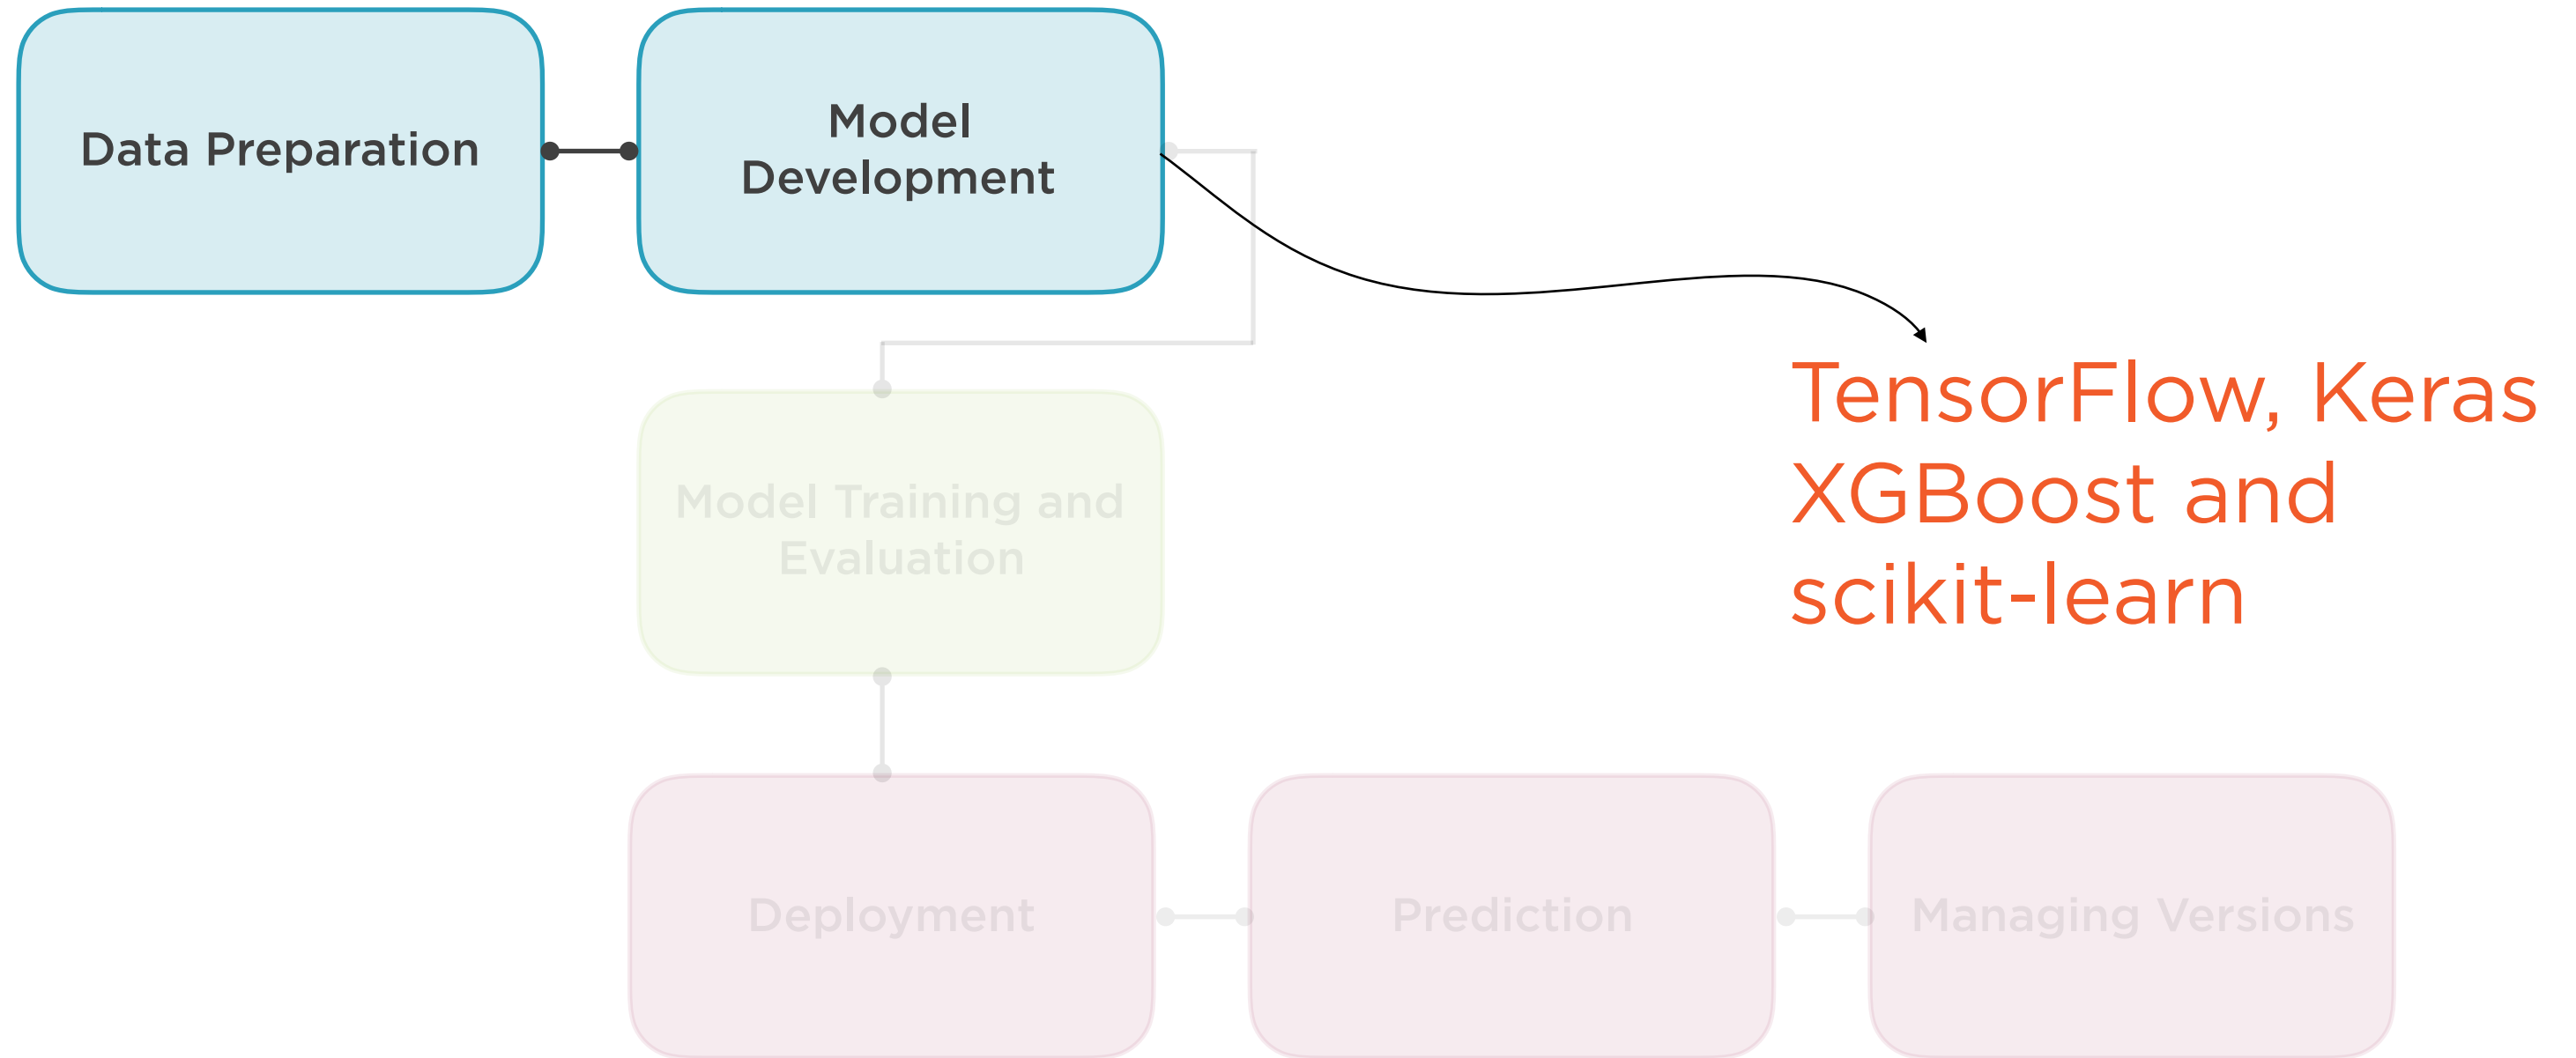
Depth of tree, minimum number of samples per node

Hyperparameter Tuning


Machine Learning Workflow



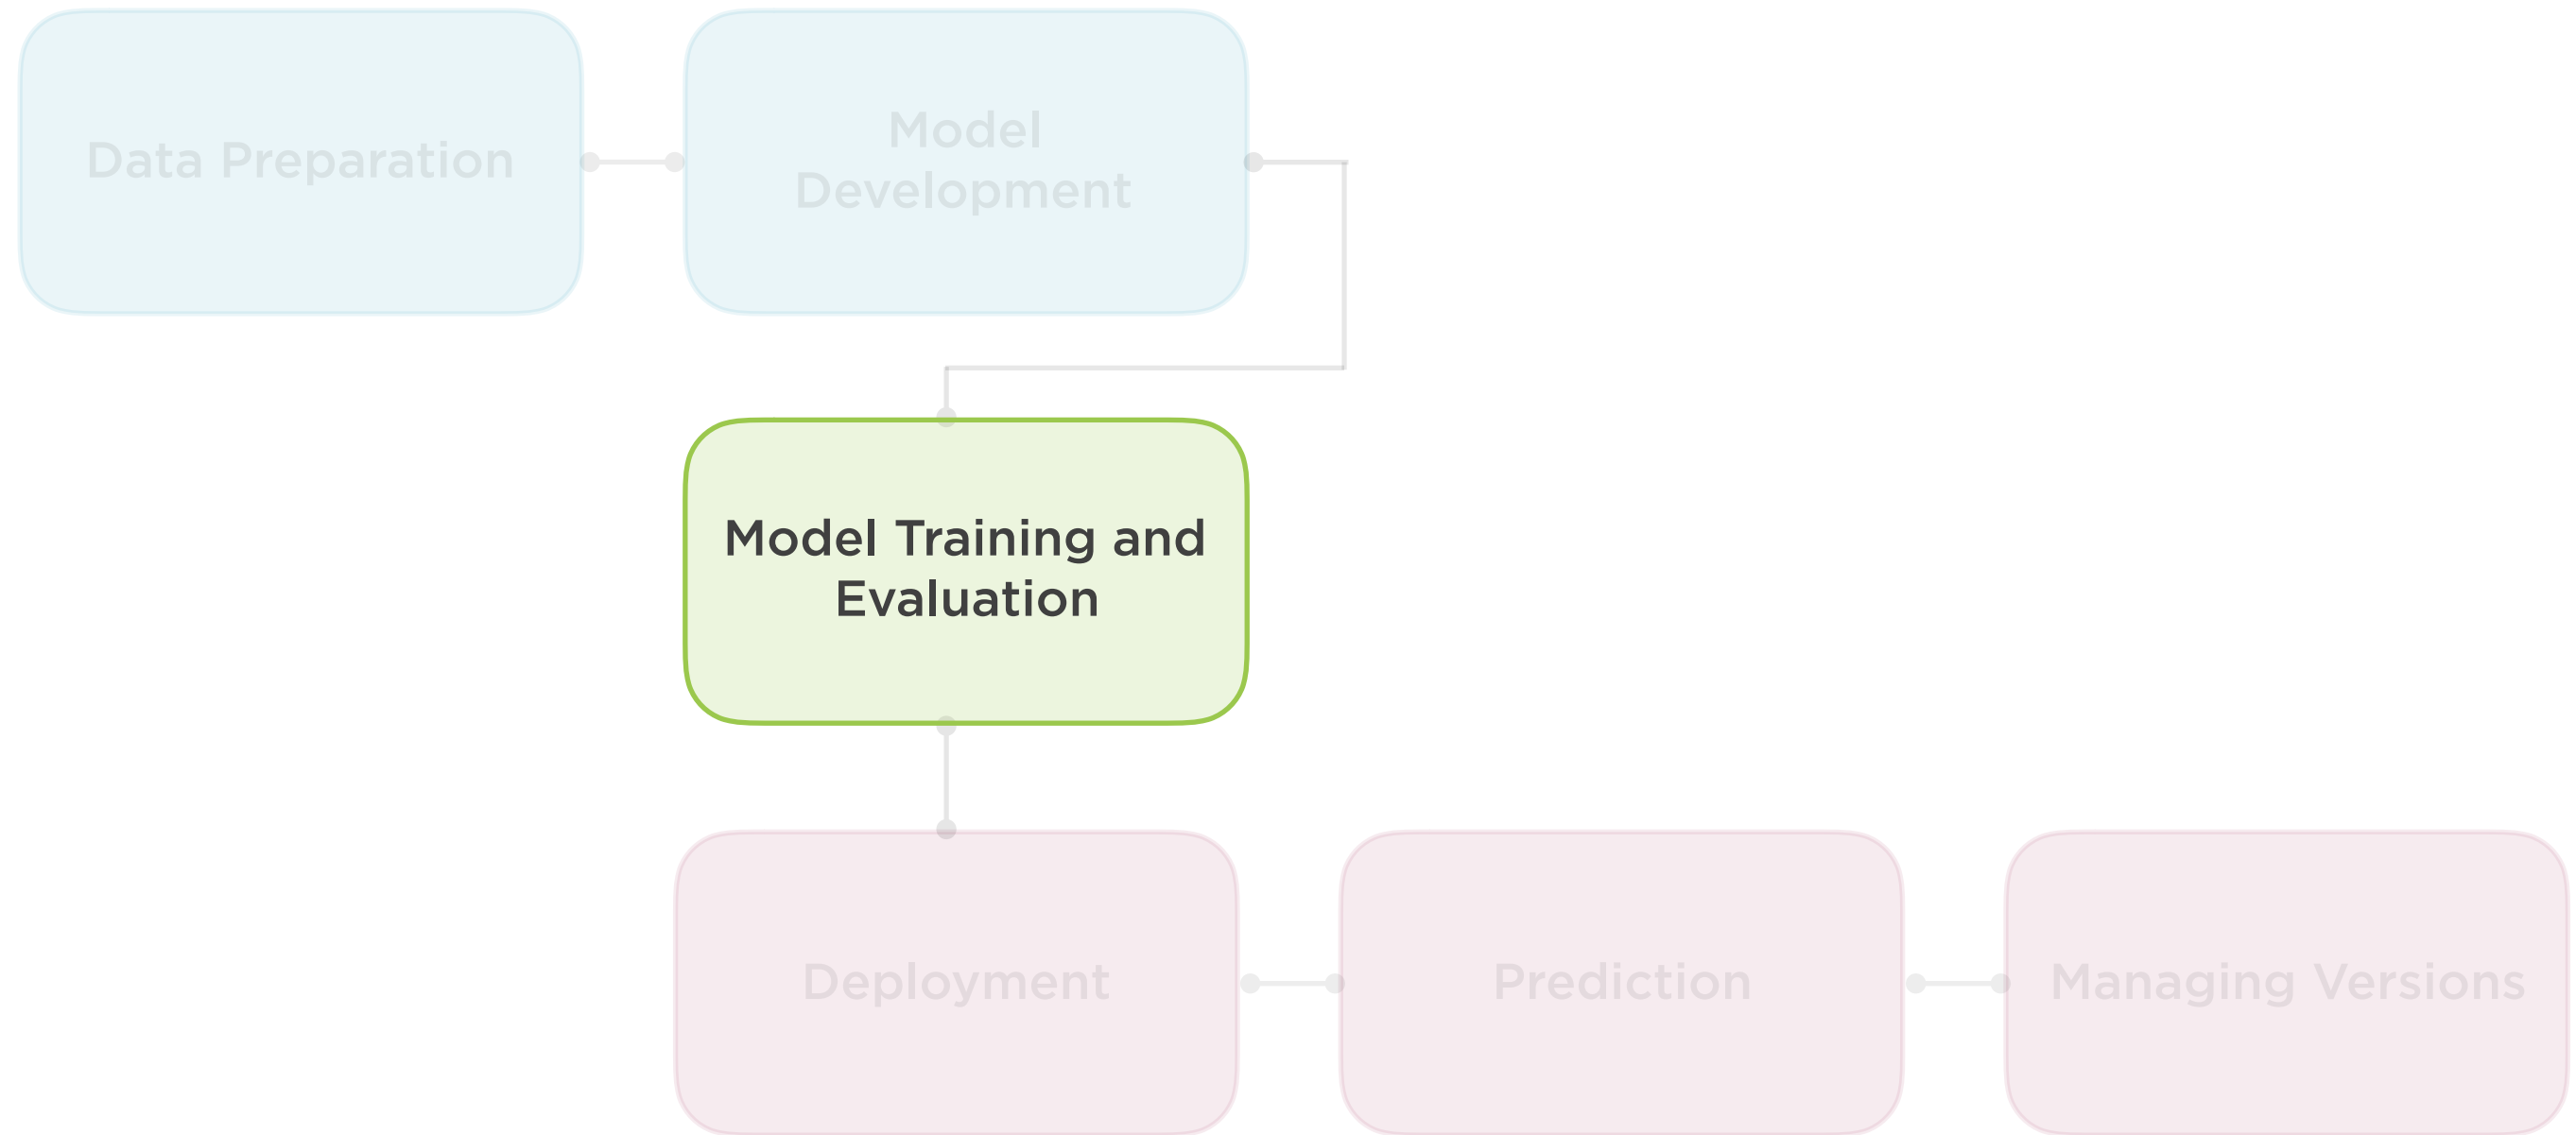
Machine Learning Workflow



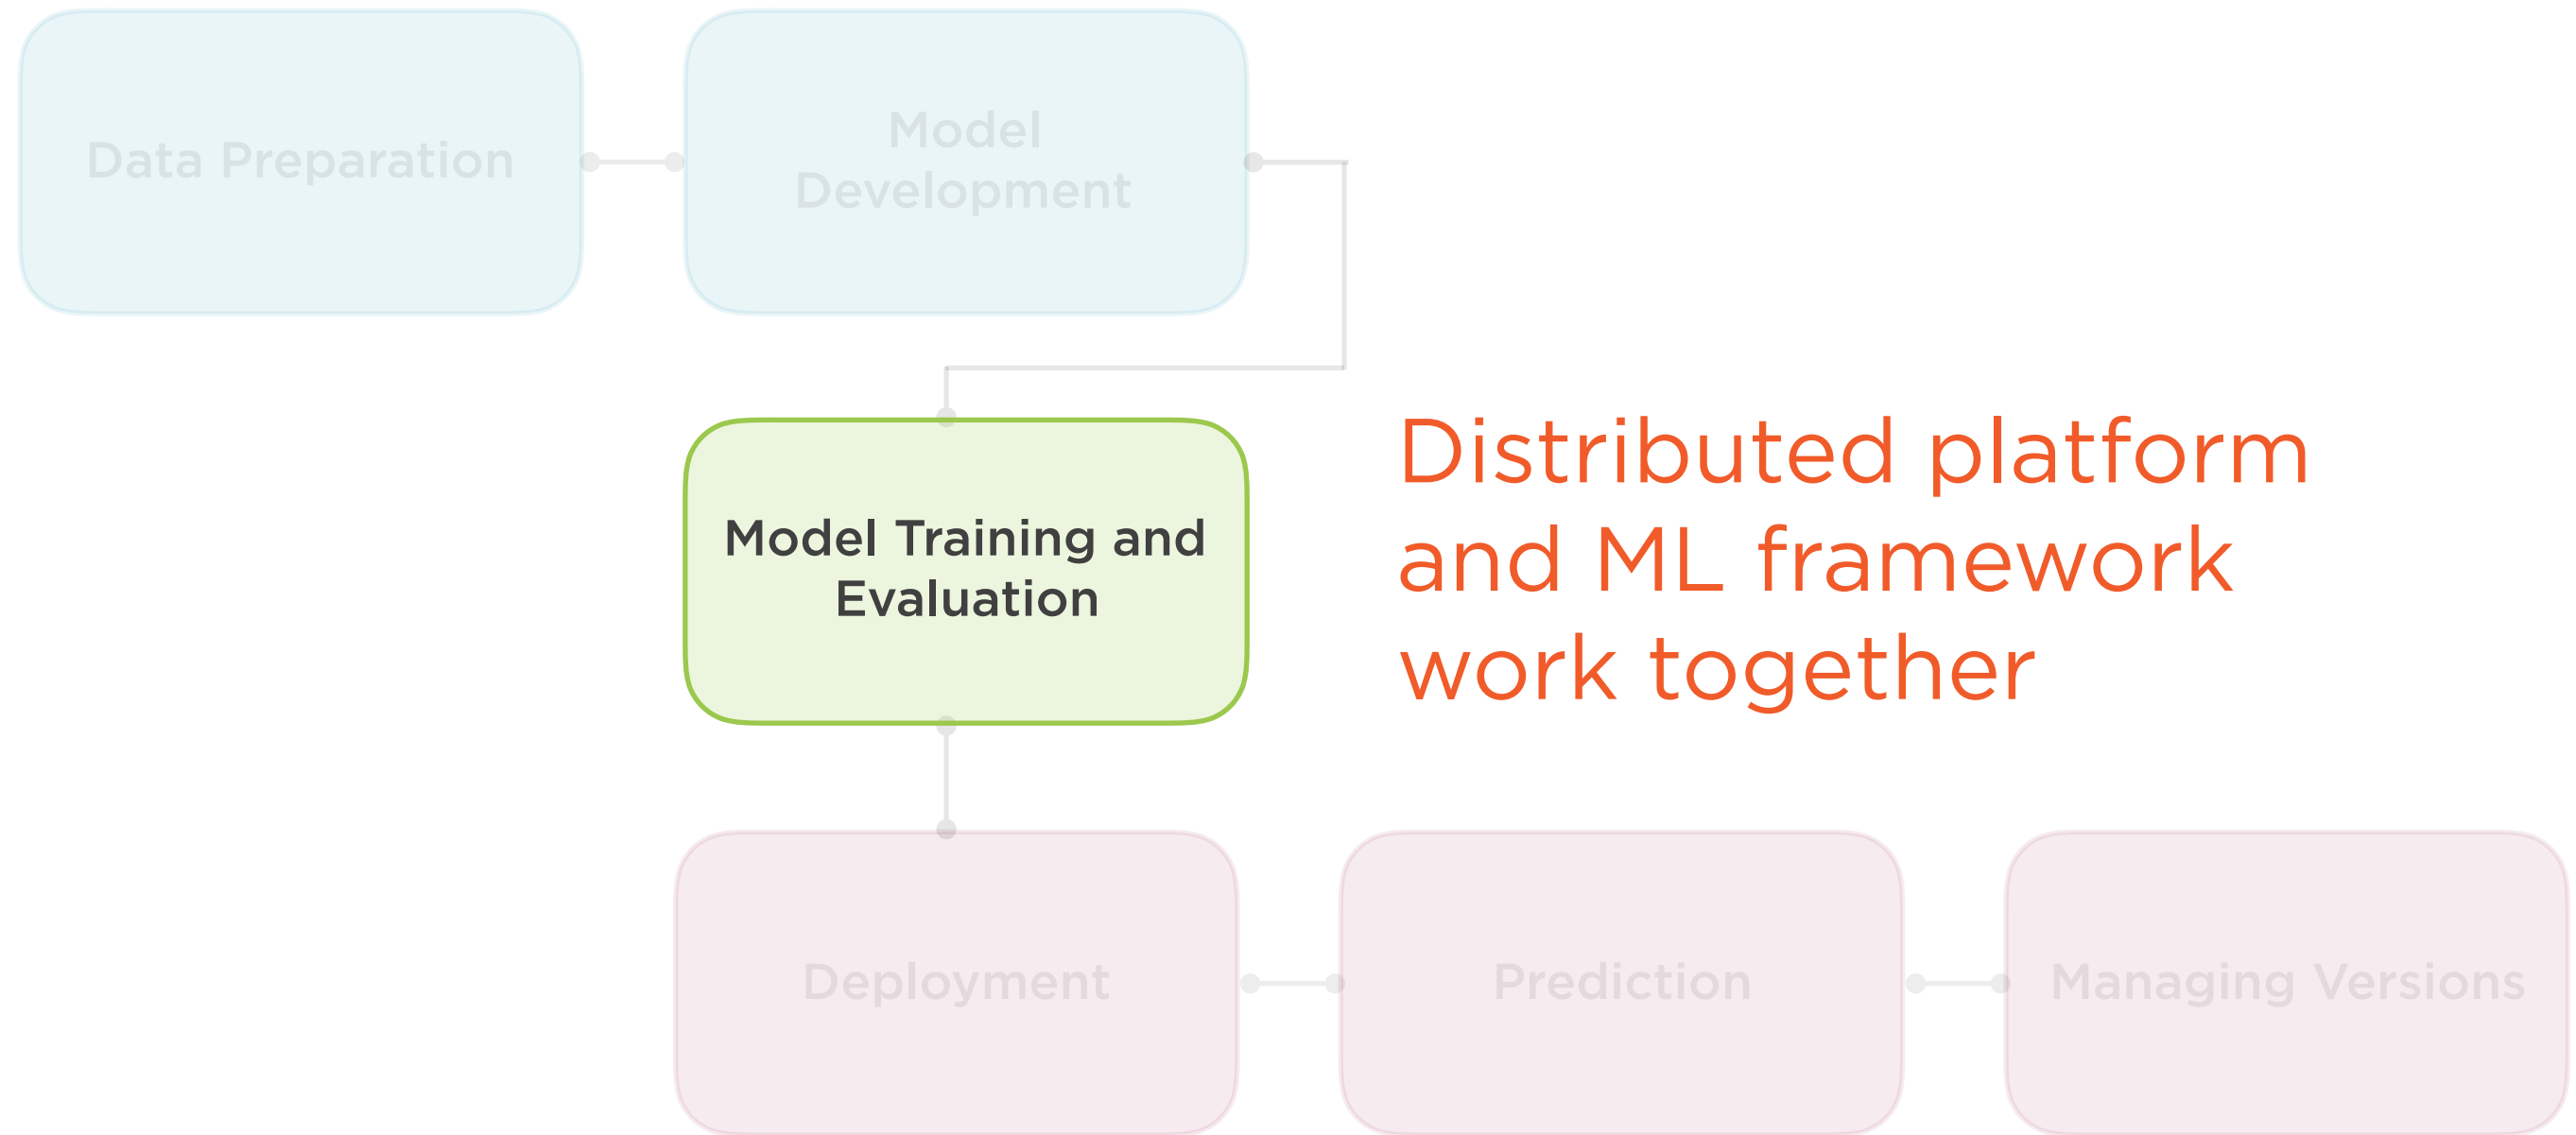
Machine Learning Workflow



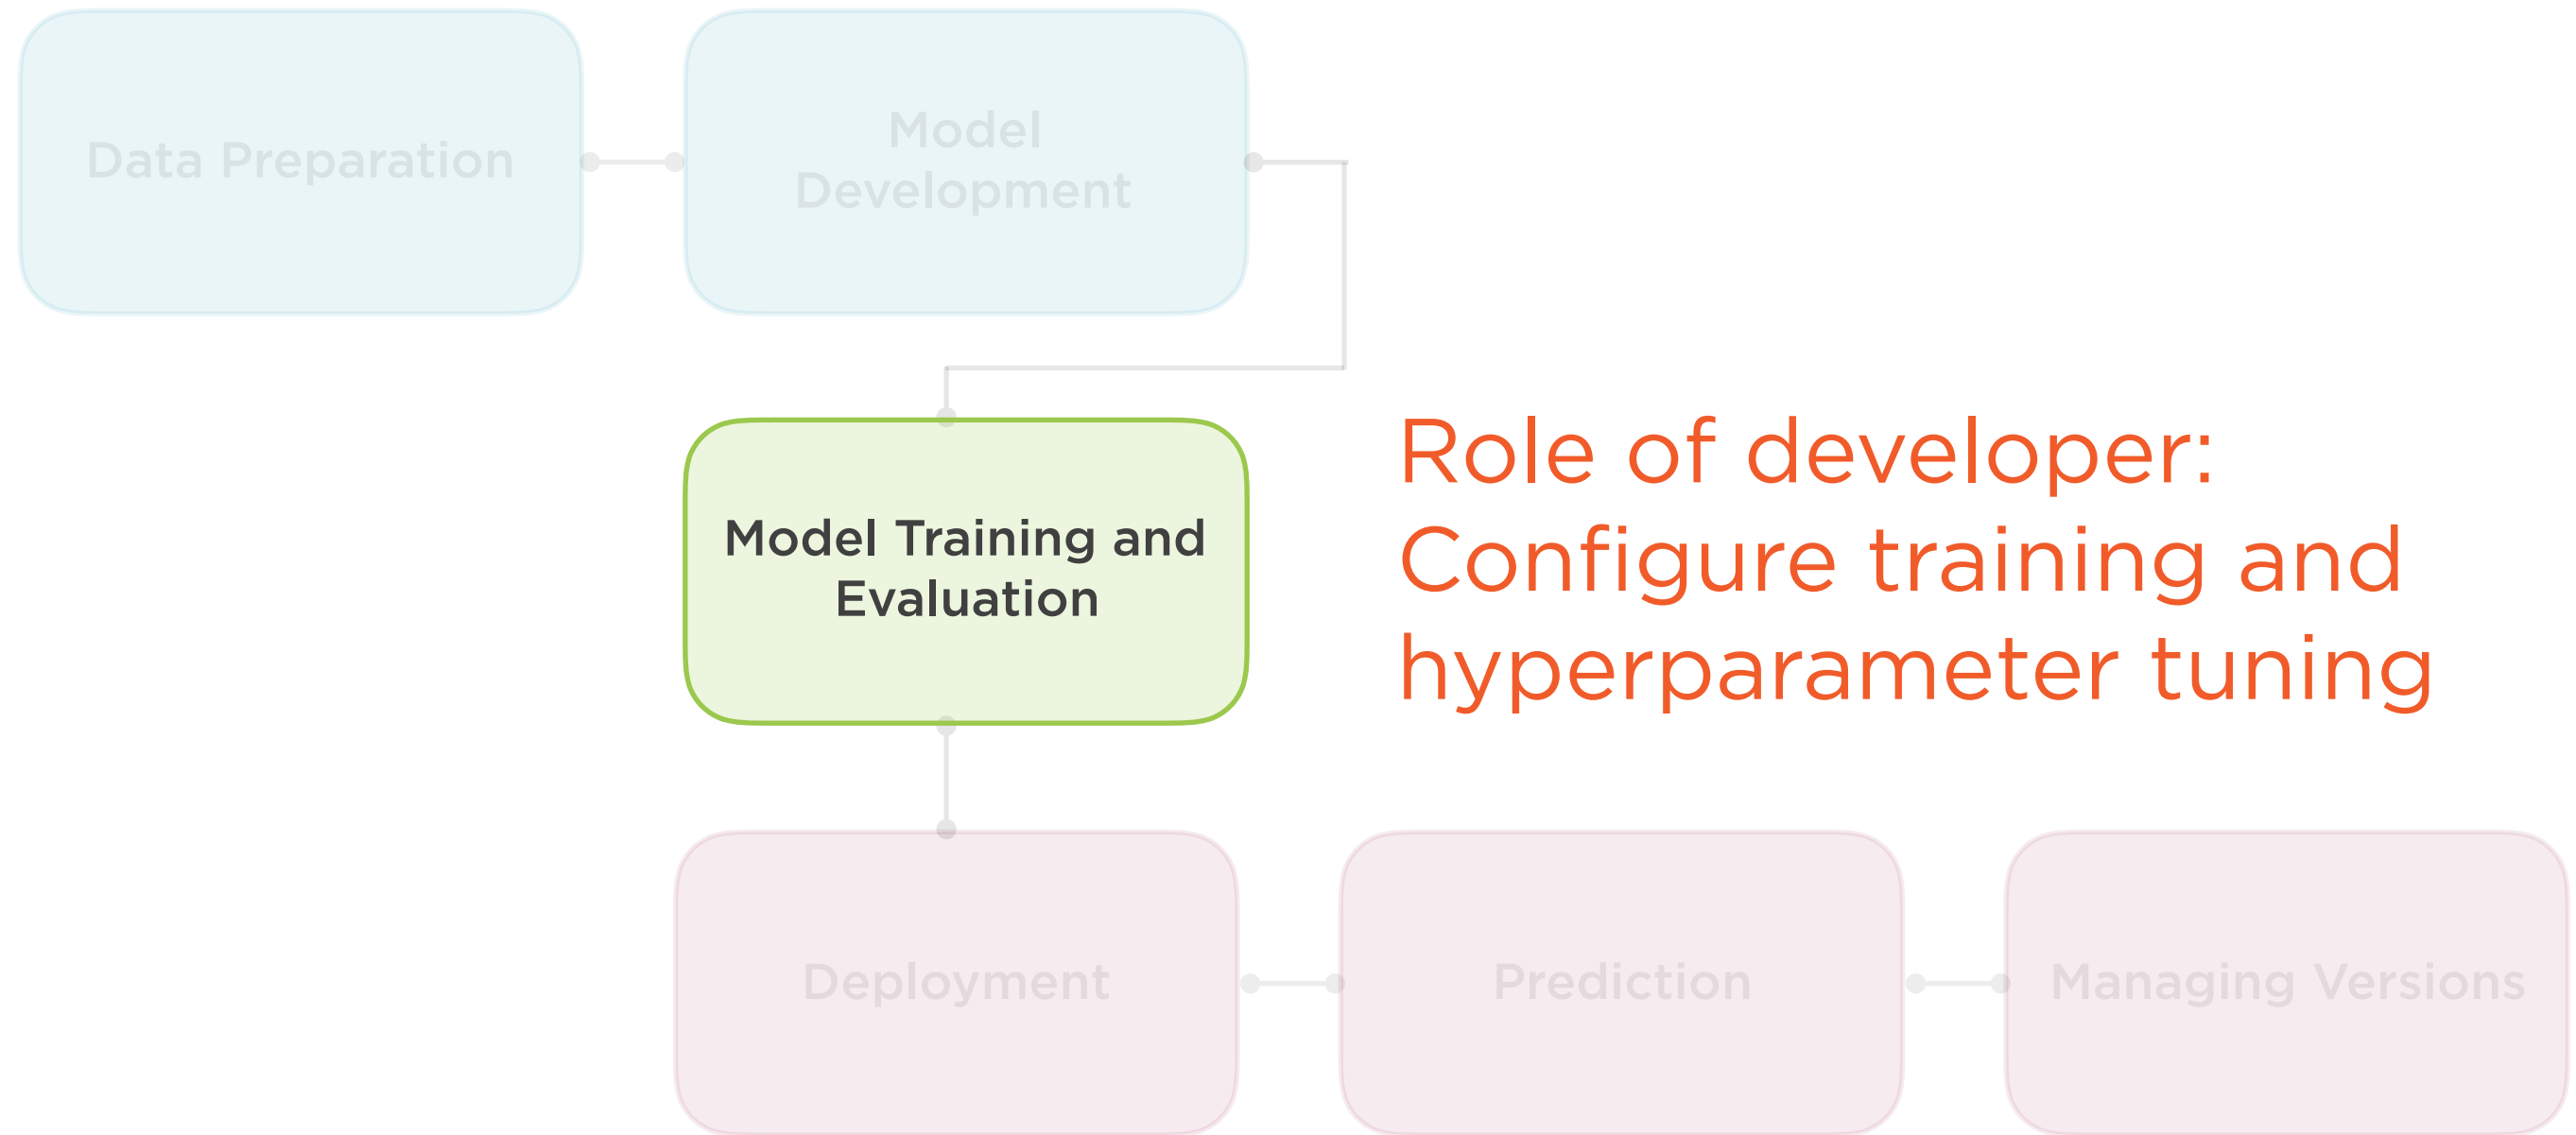
Machine Learning Workflow



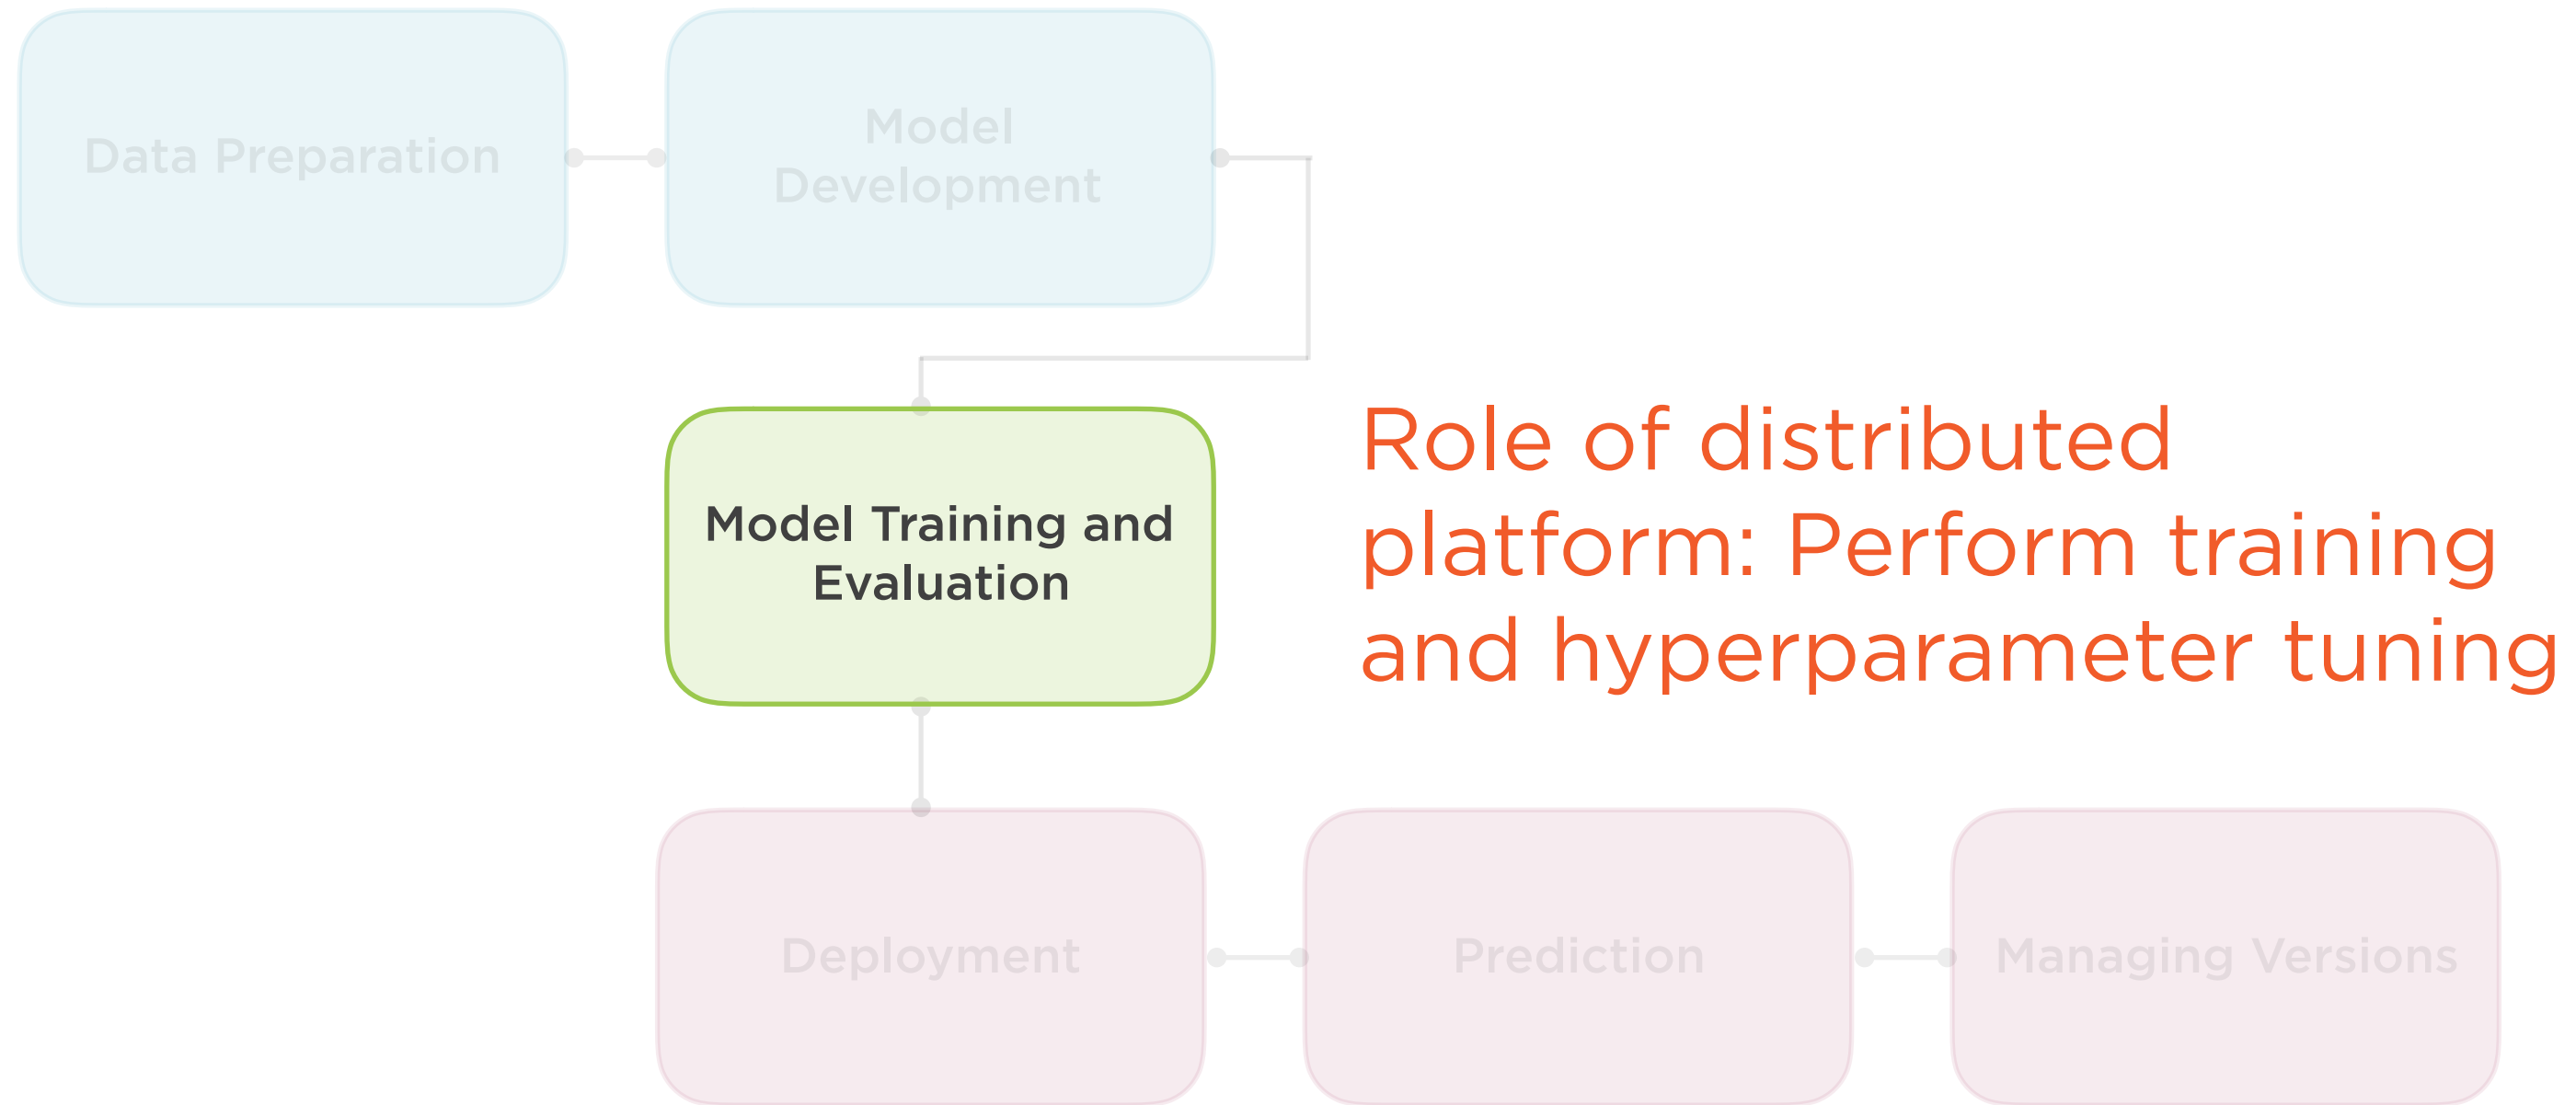
Machine Learning Workflow



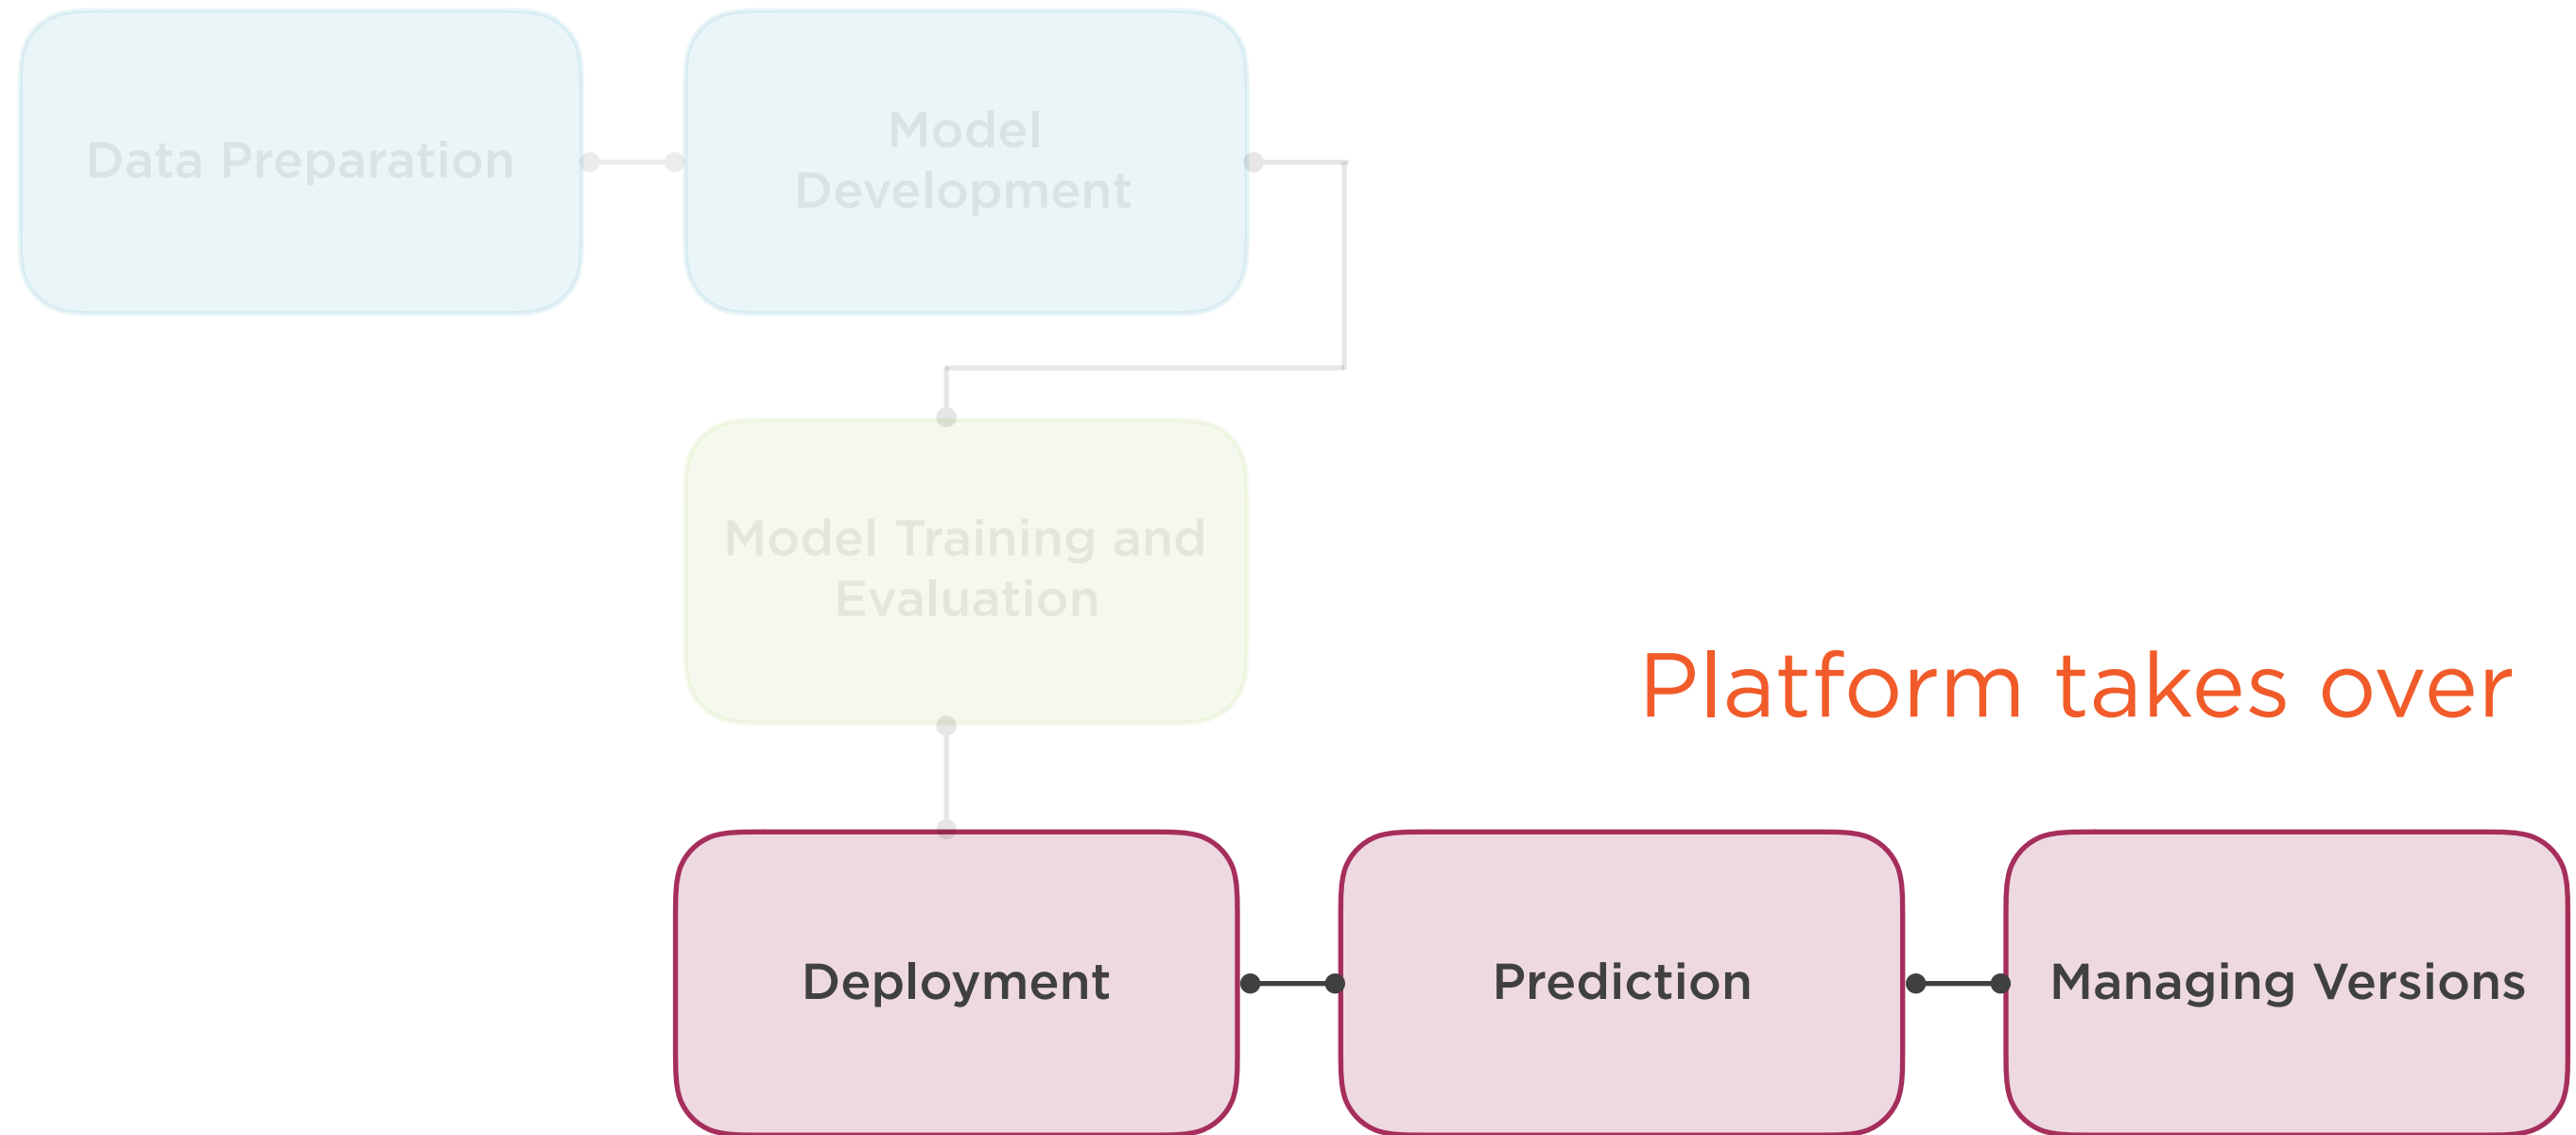
Machine Learning Workflow



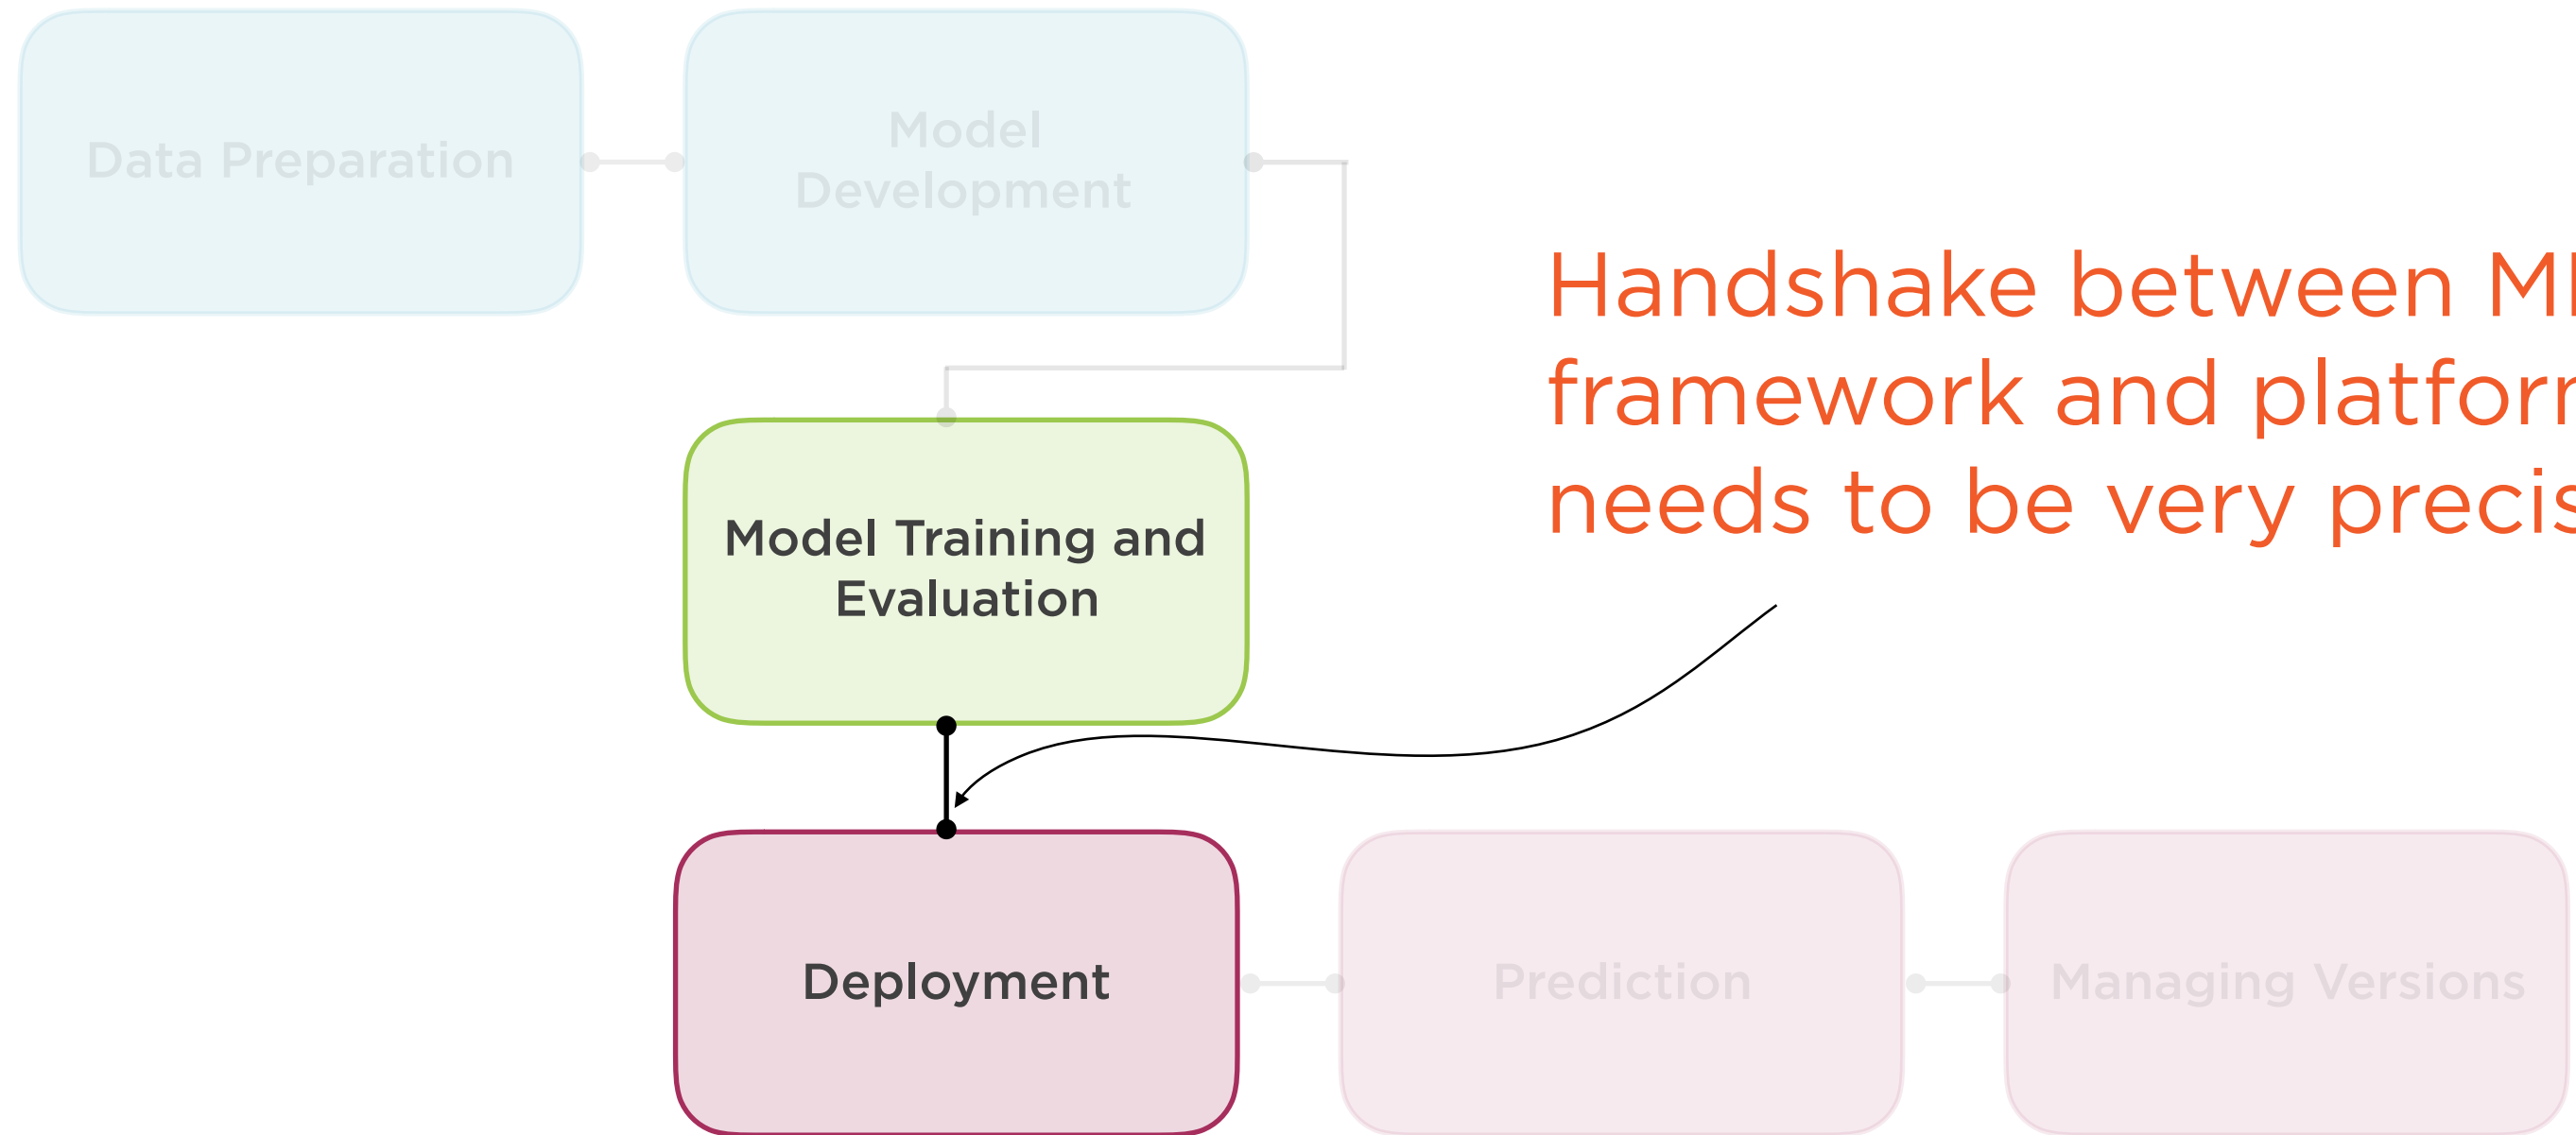
Machine Learning Workflow



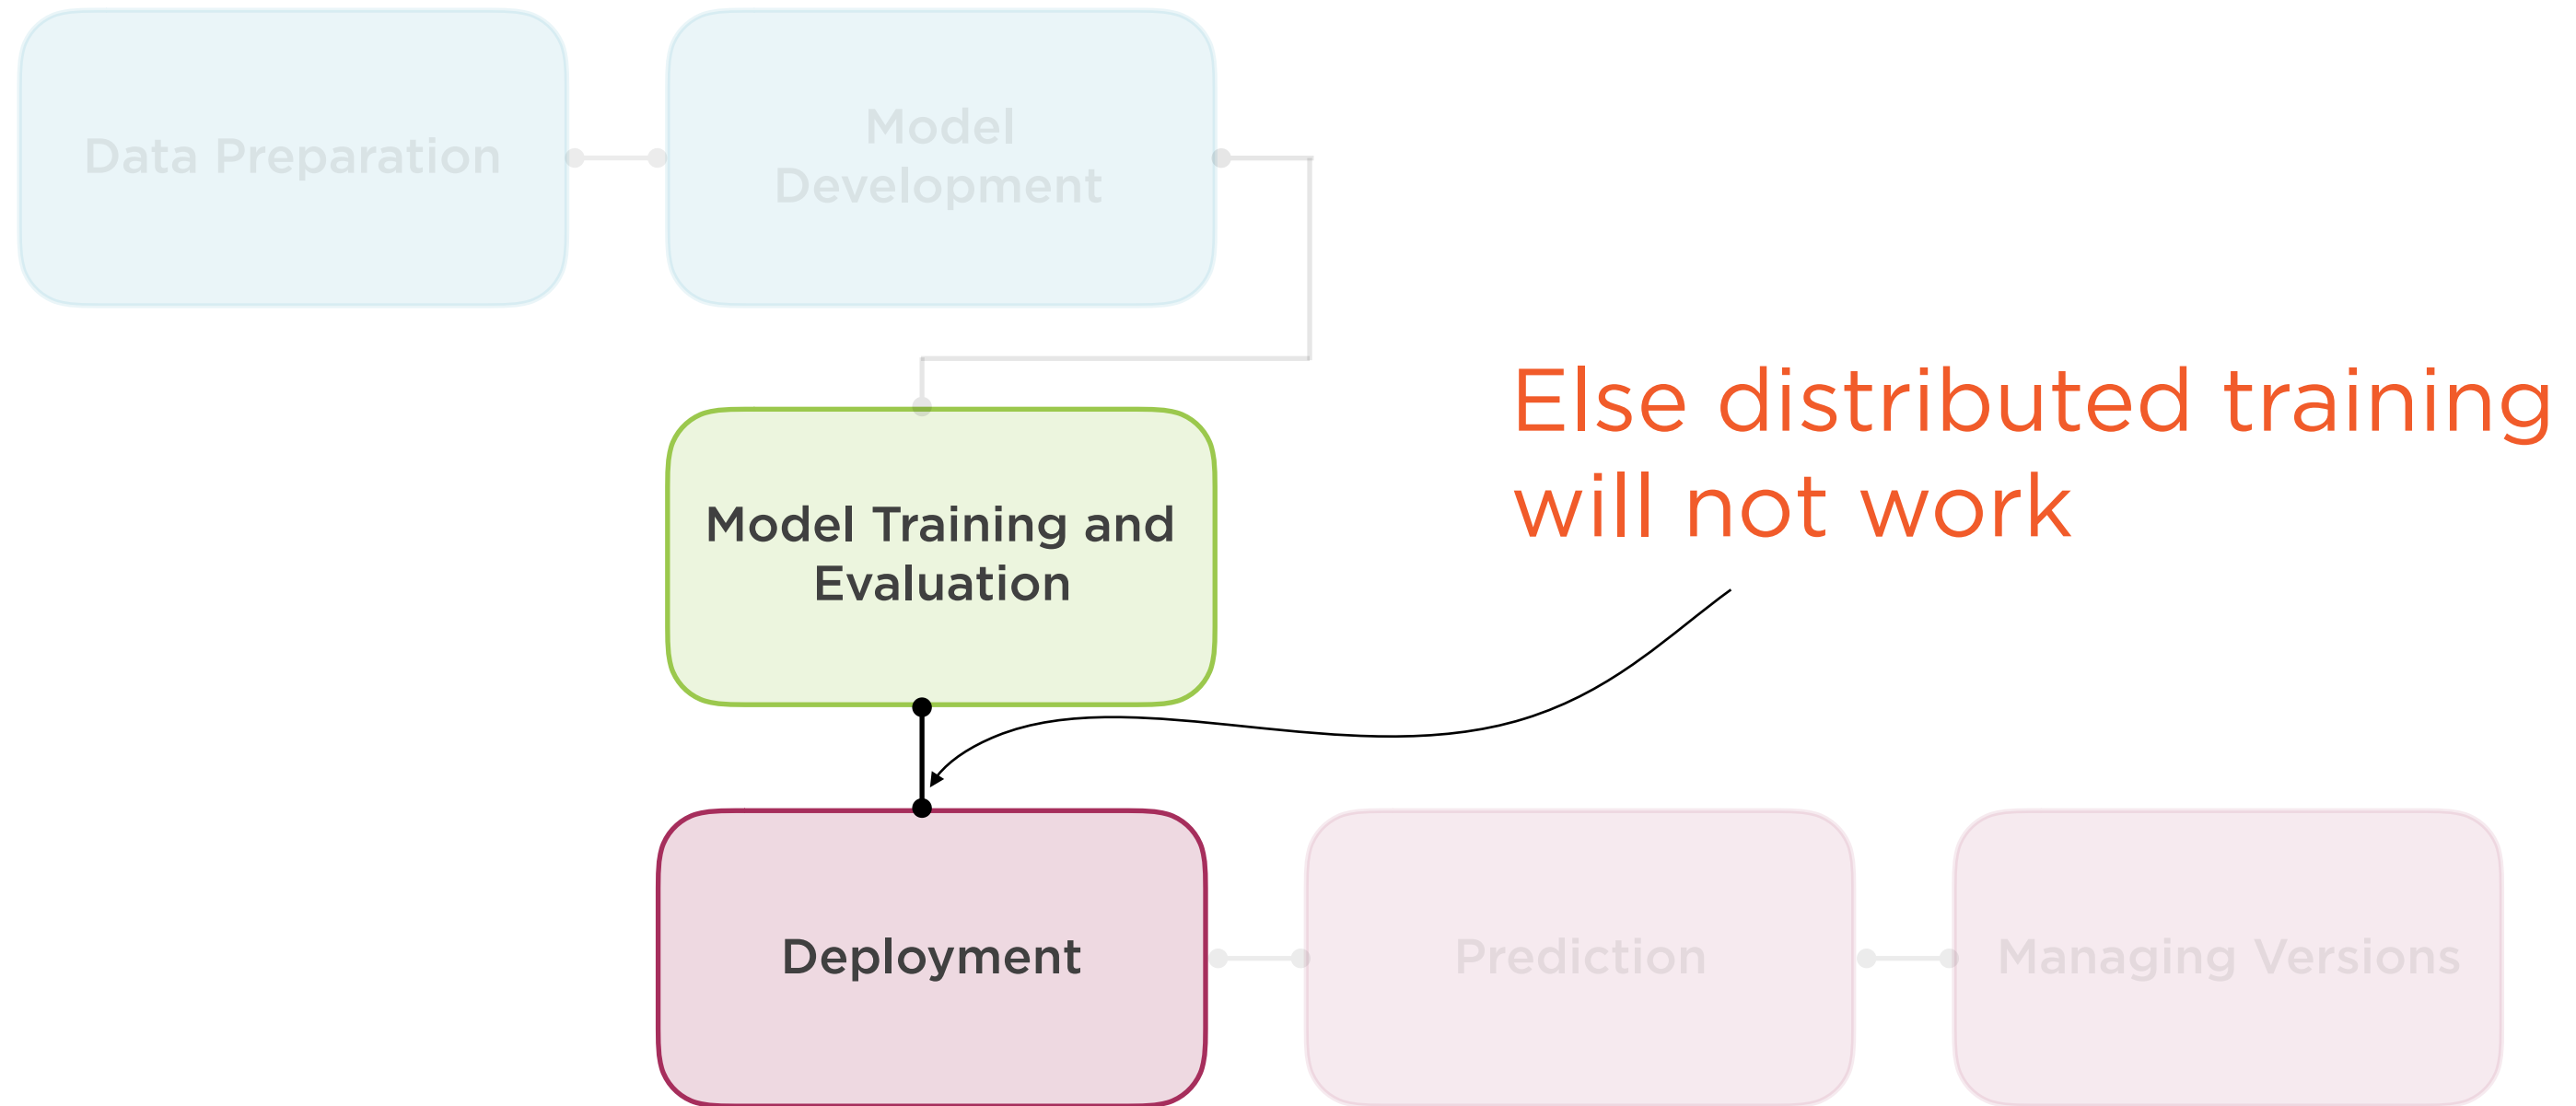
Machine Learning Workflow



Machine Learning Workflow



Machine Learning Workflow

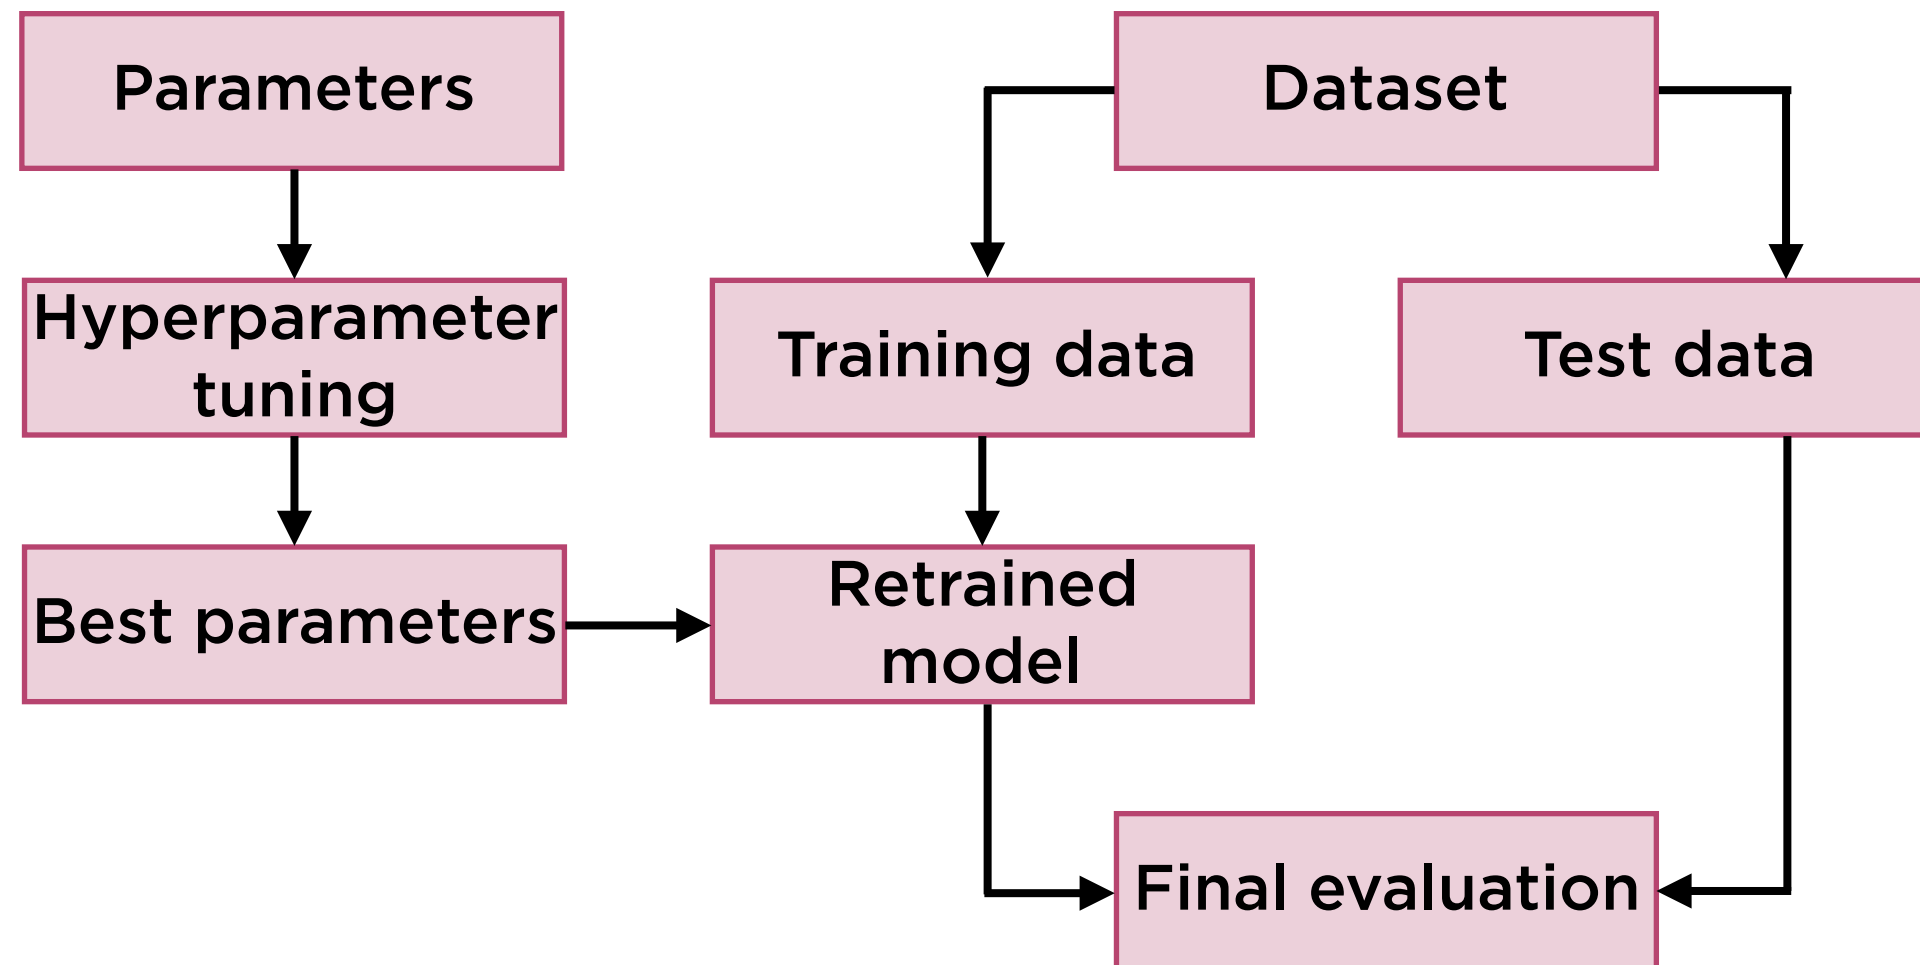


Hyperparameter tuning

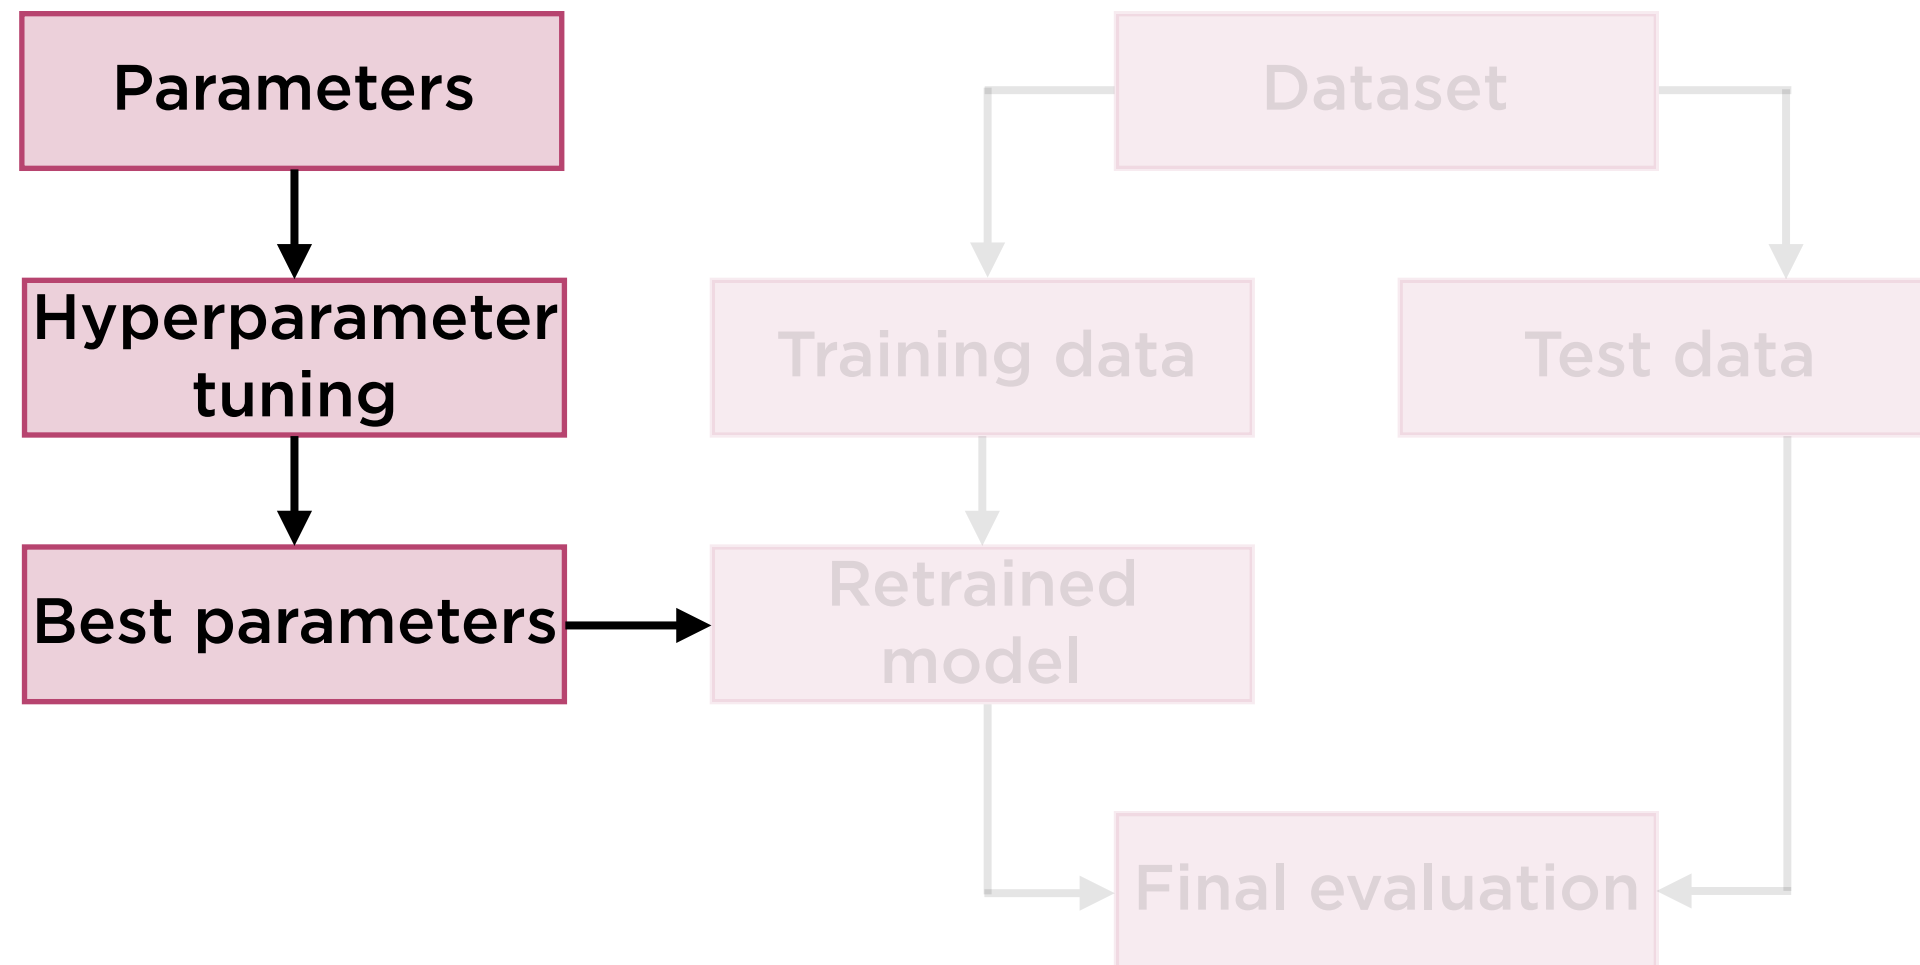
involves finding the best
design for a model

Training involves getting the
model to learn from data

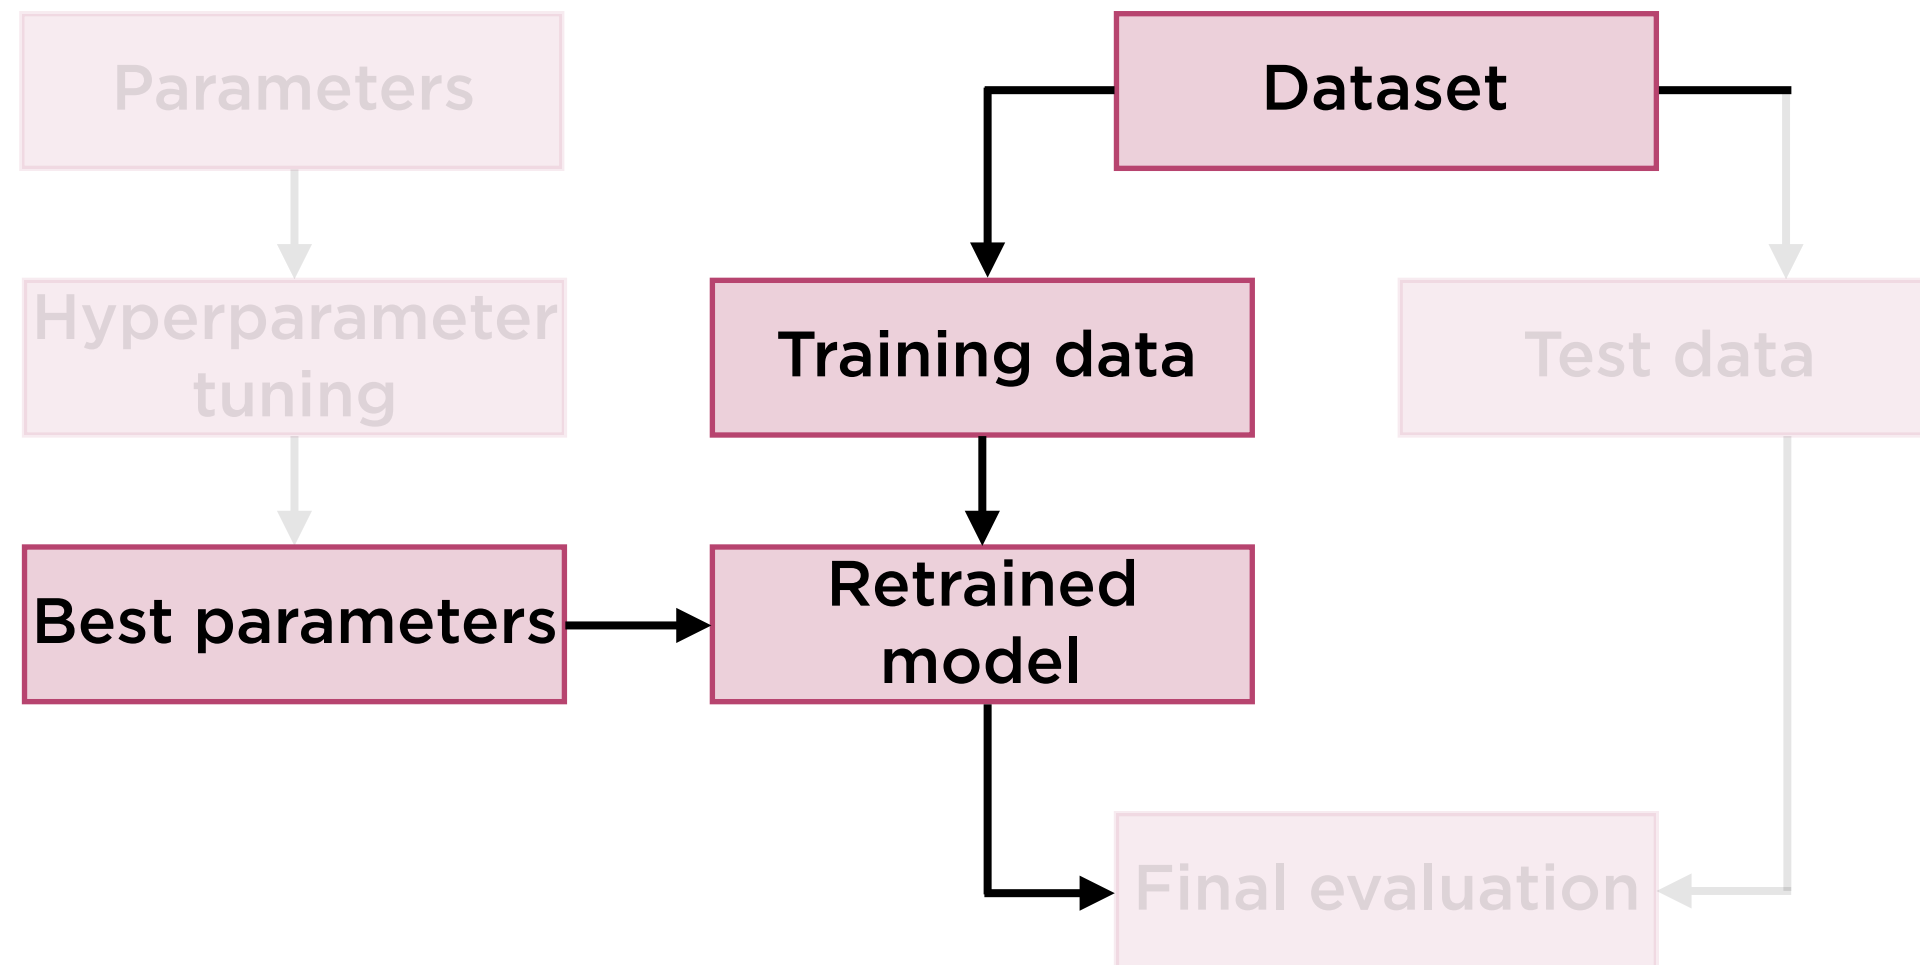
Evaluating Model Performance



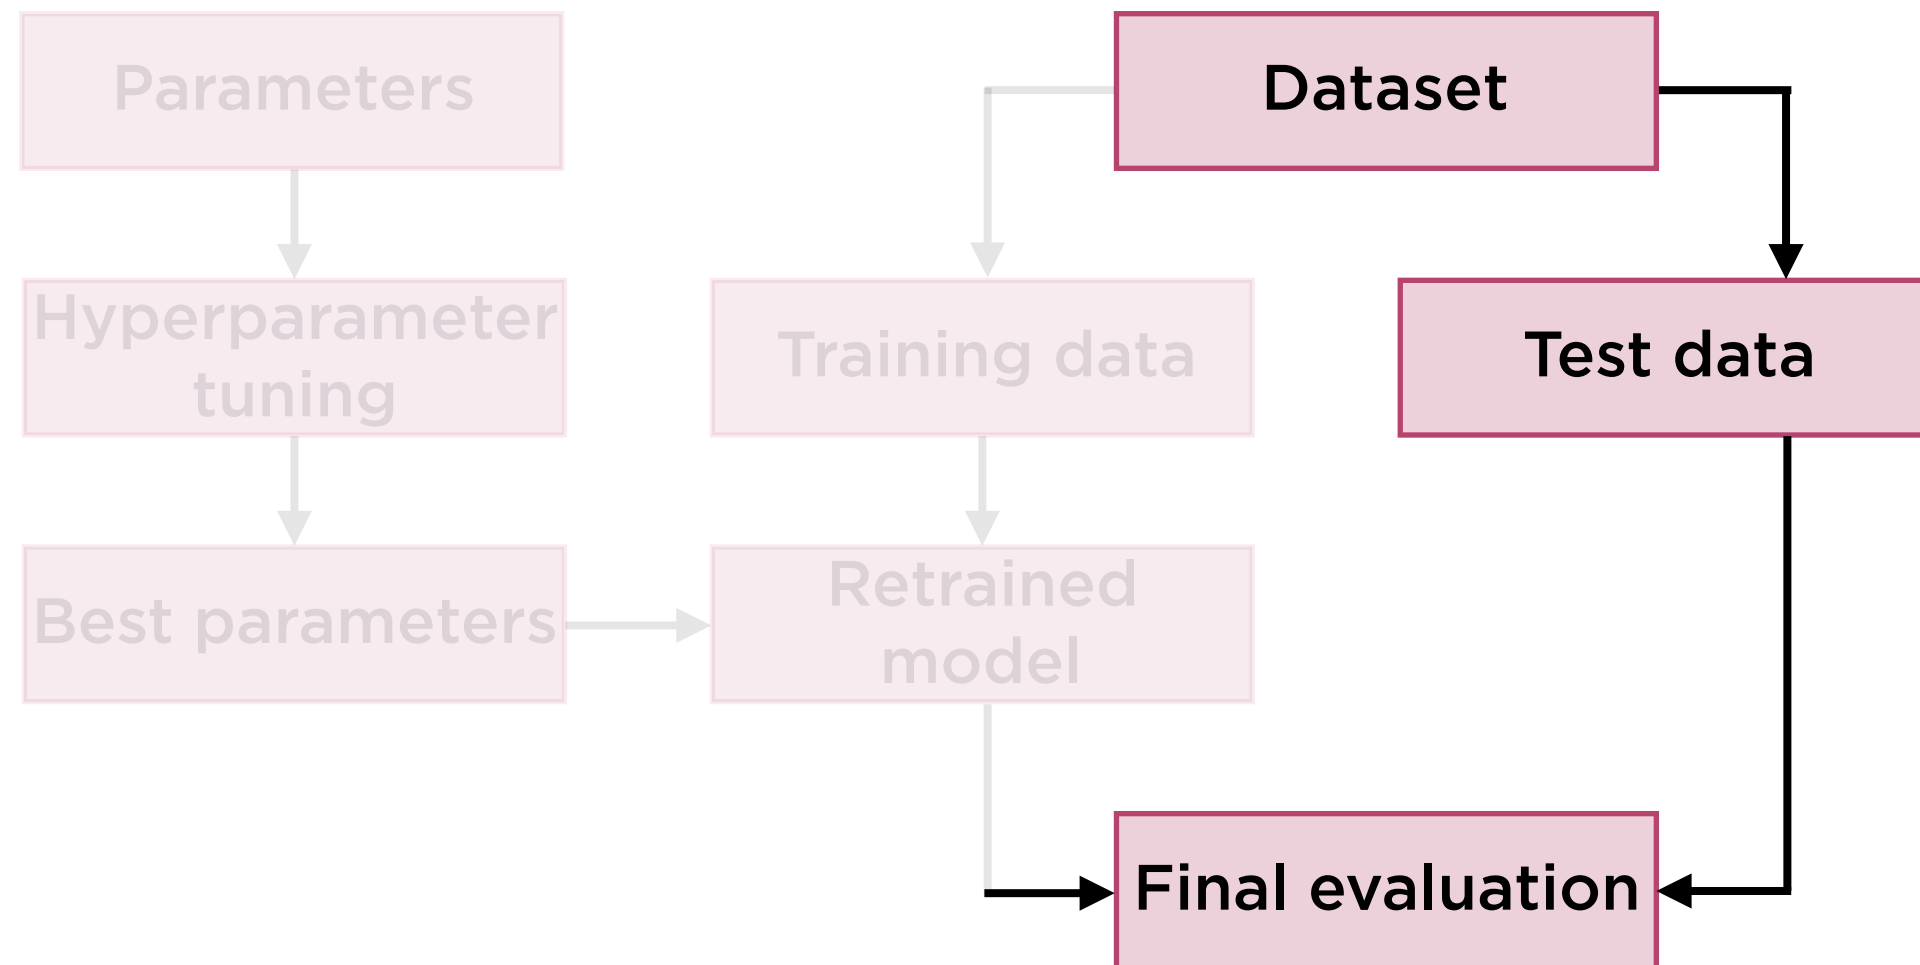
Hyperparameter Tuning to Find Best Model



Train Model with Best Design



Evaluate Model



Decision Trees

Jockey or Basketball Player?



Jockeys

Tend to be light to meet horse carrying limits



Basketball Players

Tend to be tall, strong and heavy

Jockey or Basketball Player?




Intuitively know

- jockeys tend to be light...
- ...and not very tall
- basketball players tend to be tall
- ...and also quite heavy

Fit knowledge
into rules

Each rule involves
a threshold


Decision Tree



Order of decision
variables matters

Rules and order
found using ML


Decision Tree




Classification And
Regression Tree

“CART”

Decision Tree




Decision Trees for Classification



To solve

- Traverse tree to find right node
- Return **most frequent label** of all training data points in that node

Advantages of Decision Trees




“White Box” ML ~ leverage experts

Non-parametric

- Little hyperparameter tuning
- Little data prep

Drawbacks of Decision Trees




Prone to overfitting

- Common risk with non-parametric

Unstable

- Small changes in data cause big changes in model

Random Forests



Extremely powerful technique

Example of **ensemble** learning

Individual trees should be as **different**
as possible

Hyperparameters in Decision Trees

Splitting strategy

Max depth

Min samples split

Min samples leaf

Min weight fraction

Max features

Titlecase

splitter : *string, optional (default="best")*

The strategy used to choose the split at each node. Supported strategies are "best" to choose the best split and "random" to choose the best random split.

max_depth : *int or None, optional (default=None)*

The maximum depth of the tree. If None, then nodes are expanded until all leaves are pure or until all leaves contain less than min_samples_split samples.

min_samples_split : *int, float, optional (default=2)*

The minimum number of samples required to split an internal node:

- If int, then consider `min_samples_split` as the minimum number.
- If float, then `min_samples_split` is a fraction and `ceil(min_samples_split * n_samples)` are the minimum number of samples for each split.

Changed in version 0.18: Added float values for fractions.

min_samples_leaf : *int, float, optional (default=1)*

The minimum number of samples required to be at a leaf node. A split point at any depth will only be considered if it leaves at least `min_samples_leaf` training samples in each of the left and right branches. This may have the effect of smoothing the model, especially in regression.

- If int, then consider `min_samples_leaf` as the minimum number.
- If float, then `min_samples_leaf` is a fraction and `ceil(min_samples_leaf * n_samples)` are the minimum number of samples for each node.

Changed in version 0.18: Added float values for fractions.

min_weight_fraction_leaf : *float, optional (default=0.)*

The minimum weighted fraction of the sum total of weights (of all the input samples) required to be at a leaf node. Samples have equal weight when sample_weight is not provided.

max_features : *int, float, string or None, optional (default=None)*

The number of features to consider when looking for the best split:

- If int, then consider `max_features` features at each split.
- If float, then `max_features` is a fraction and `int(max_features * n_features)` features are considered at each split.
- If "auto", then `max_features=sqrt(n_features)`.
- If "sqrt", then `max_features=sqrt(n_features)`.
- If "log2", then `max_features=log2(n_features)`.
- If None, then `max_features=n_features`.

Note: the search for a split does not stop until at least one valid partition of the node samples is found, even if it requires to effectively inspect more than `max_features` features.

random_state : *int, RandomState instance or None, optional (default=None)*

If int, random_state is the seed used by the random number generator; If RandomState instance, random_state is the random number generator; If None, the random number generator is the RandomState instance used by `np.random`.

max_leaf_nodes : *int or None, optional (default=None)*

Grow a tree with `max_leaf_nodes` in best-first fashion. Best nodes are defined as relative reduction in impurity. If None then unlimited number of leaf nodes.

min_impurity_decrease : *float, optional (default=0.)*

A node will be split if this split induces a decrease of the impurity greater than or equal to this value.

The weighted impurity decrease equation is the following:

$$N_t / N * (impurity - N_{t_R} / N_t * right_impurity - N_{t_L} / N_t * left_impurity)$$

where `N` is the total number of samples, `Nt` is the number of samples at the current node, `Nt_L` is the number of samples in the left child, and `Nt_R` is the number of samples in the right child.

`N`, `Nt`, `Nt_R` and `Nt_L` all refer to the weighted sum, if `sample_weight` is passed.

New in version 0.19.

min_impurity_split : *float, (default=1e-7)*

Threshold for early stopping in tree growth. A node will split if its impurity is above the threshold, otherwise it is a leaf.

Deprecated since version 0.19: `min_impurity_split` has been deprecated in favor of `min_impurity_decrease` in 0.19. The default value of `min_impurity_split` will change from 1e-7 to 0 in 0.23 and it will be removed in 0.25. Use `min_impurity_decrease` instead.

Demo

**Hyperparameter tuning using Azure
ML Studio**

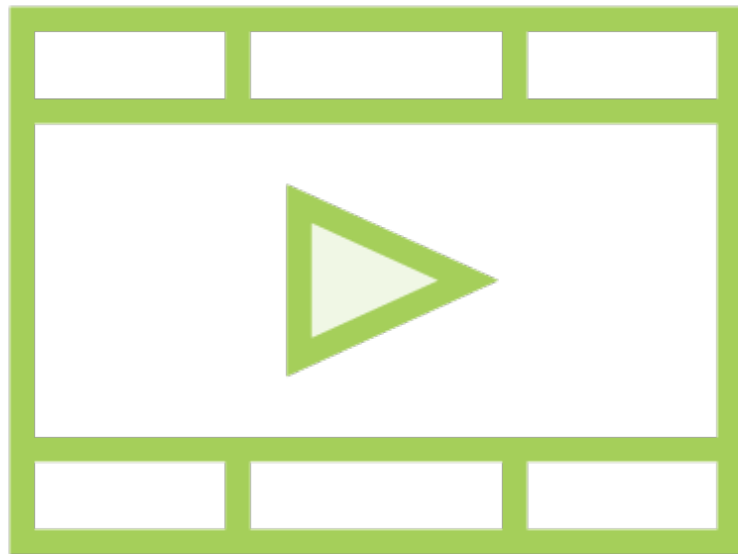
Summary

Hyperparameter tuning in model validation

Model parameters vs. model hyperparameters

Hyperparameter tuning using Azure ML Studio

Related Courses



Summarizing Data and Deducing Probabilities

Communicating Data Insights